CS 760: Machine Learning Spring 2024

Homework 5: Nearest Neighbors & Naive Bayes

AUTHORS: Jed Pulley

DO NOT POLLUTE! AVOID PRINTING, OR PRINT 2-SIDED MULTIPAGE.

Problem 5.1

(a) See code in *KNN.ipynb* under Problem 5.1. I chose to implement K-Nearest Neighbors where $K = 5$. I went this route since it's as simple as normal Nearest Neighbors to implement, but more robust.

(b) I used *np.linalg.norm* to implement the euclidean distance. I chose the L2 norm because I find it to be more straightforward and intuitive, plus I'm much more used to using it from my previous Linear Algebra classes.

(c) See code in *KNN.ipynb* under Problem 5.1. I tested out multiple features and summed up the counts of survival, based on how many K nearest neighbors I use. Unfortunately, I did not survive based on my demographics. Notably, as I increased K, the number of survivers went down.

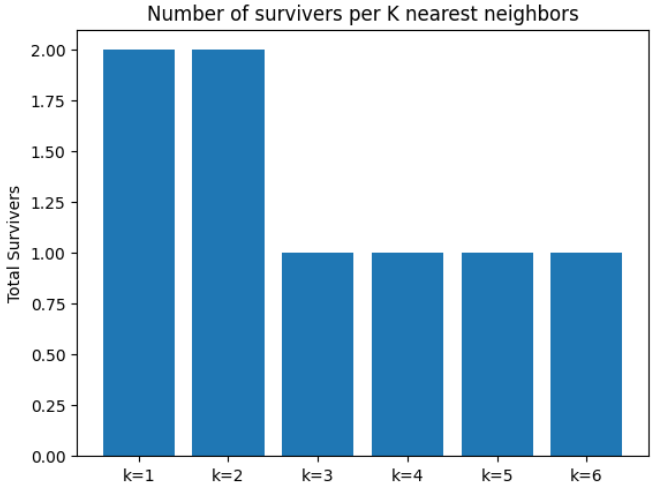

(d) While $k = 1$ and $k = 2$ have the most survivors among my samples, I beleive that's because there isn't enough wiggle room for correct classification. I believe $k = 6$ is most representative give the fact that I have 6 different features.

(e) The most apparent solution is to run multiple rounds of cross validation along different values for K to assess the accuracy.

Problem 5.2

(a)

(b)

(c)

(d)

Problem 5.3

Problem 5.4

Problem 5.5